

## Project 4 : Stack Overflow - askubuntu dataset










Adarsh and Ayush

The project involves extracting data from several smaller datasets and combining them together to do analysis.

### Dataset:

The dataset contains logs for askubuntu stackexchange logs (<https://askubuntu.com>) in XML format.

- The total size of the dataset is 22 GB.
- It is stored in the GCS bucket <gs://stackoverflow-dataset-677>
- The dataset includes multiple xml files corresponding to different attributes of the dataset.

|                          |   |
|--------------------------|---|
| <input type="checkbox"/> |  Badges.xml      |
| <input type="checkbox"/> |  Comments.xml    |
| <input type="checkbox"/> |  Comments_small  |
| <input type="checkbox"/> |  PostHistory.xml |
| <input type="checkbox"/> |  PostLinks.xml   |
| <input type="checkbox"/> |  Posts.xml     |
| <input type="checkbox"/> |  Tags.xml      |
| <input type="checkbox"/> |  Users.xml     |
| <input type="checkbox"/> |  Votes.xml     |

- The following are the relevant features for each XML file.
  - **Users:** Reputation, CreationDate, DisplayName, WebsiteUrl, Location, Views, UpVotes, DownVotes, AccountId
  - **Posts:** Id, PostTypeId, AcceptedAnswerId, CreationDate, Score, Body, OwnerUserId, Title, Tags, AnswerCount, CommentCount
  - **Comments :** RowId, PostId, Score, Text, CreationDate, UserId, ContentLicense
  - **Tags:** RowId, TagName, Count, ExcerptPostId, WikiPostId
  - **Badges :** RowId, UserId, Name, Date, Class, TagBased
  - **Votes :** RowId, PostId, VoteTypeId, CreationDate

### Tasks to be completed:

1. a) Extracting user id and username and storing it to GCS.  
b) Extracting the comments and doing inner join with Spark SQL and display to user.

2. Trending Users with max comments: Finding users who posted maximum comments and visualizing it.
3. Trending Topics in Comments: Extracting the trending topics in comments related to ubuntu and visualizing it.
4. Automating task 1 with Apache Airflow i.e. in a single pipeline

## Server Setup:

We used Google Cloud DataProc to setup a pyspark cluster of 1 Master Nodes (2 vCPUs ) and 6 Worker Nodes (1vCPU) each with spark 3.2, python 3 and ubuntu 18.04 operating system.

For first 3 tasks, the server setup was done manually.

## Manual Setup:

Google Cloud Platform big-data-final Search products and resources

Dataproc Create a cluster

Clusters

- Set up cluster  
Begin by providing basic information.
- Configure nodes (optional)  
Change node compute and storage capabilities.
- Customize cluster (optional)  
Add cluster properties, features, and actions.
- Manage security (optional)  
Change access, encryption, and security settings.

CREATE CANCEL

Equivalent [REST](#) or [command line](#)

**Name**  
Cluster Name  
pyspark-stack-overflow

**Location**  
Region  
us-west2  
Zone  
us-west2-b

**Cluster type**  
☒ Standard (1 master, N workers)  
☐ Single Node (1 master, 0 workers)  
Provides one node that acts as both master and worker. Good for proof-of-concept or small-scale processing  
☐ High Availability (3 masters, N workers)  
Hadoop High Availability mode provides uninterrupted YARN and HDFS operations despite single-node failures or reboots

**Autoscaling**  
Automates cluster resource management based on an autoscaling policy.  
Policy  
None

**Versioning**  
Use a custom image to load pre-installed packages. [Learn more](#)  
**Image Type and Version**  
preview-ubuntu18  
**Release Date**  
Preview released on 6/10/2020.  
[CHANGE](#)

Google Cloud Platform

big-data-final

Search products and resources

Dataproc

Clusters

Jobs

Workflows

Autoscaling policies

Component exchange

Notebooks

Create a cluster

Set up cluster

Begin by providing basic information.

Configure nodes (optional)

Change node compute and storage capabilities.

Customize cluster (optional)

Add cluster properties, features, and actions.

Manage security (optional)

Change access, encryption, and security settings.

CREATE

CANCEL

Equivalent [REST](#) or [command line](#)

Master node

Contains the YARN Resource Manager, HDFS NameNode, and all job drivers.

Machine family

GENERAL-PURPOSE

COMPUTE-OPTIMIZED

MEMORY-OPTIMIZED

Machine types for common workloads, optimized for cost and flexibility

Series

N1

Powered by Intel Skylake CPU platform or one of its predecessors

Machine type

n1-standard-2 (2 vCPU, 7.5 GB memory)

vCPU

2

Memory

7.5 GB

CPU PLATFORM AND GPU

Primary disk size (min 10GB)

10

GB

?

Primary disk type

Standard Persistent Disk

?

Number of local SSDs \*

0

x 375GB

?

Worker nodes

Each contains a YARN NodeManager and a HDFS DataNode. HDFS replication factor is 2.

⚠

Creating a Dataproc cluster using the n1-standard-1 machine type is not recommended. Consider using a machine type with higher memory. [Learn more](#)

Machine family

GENERAL-PURPOSE

COMPUTE-OPTIMIZED

MEMORY-OPTIMIZED

Machine types for common workloads, optimized for cost and flexibility

Series

N1

https://console.cloud.google.com/home?project=big-data-final-296808

# Result:

- Cluster

Google Cloud Platform

big-data-final

Search products and resources

Dataproc

Clusters

Jobs

Workflows

Autoscaling policies

Component exchange

Notebooks

CREATE CLUSTER

REFRESH

DELETE

REGIONS

Search clusters, press Enter

|  | Name                           | Region | Zone       | Total worker nodes | Scheduled deletion | Cloud Storage staging bucket                    | Created                  | Status       |
|--|--------------------------------|--------|------------|--------------------|--------------------|---|--------------------------|--------------|
|  | dataproc-cluster-demo-20201210 | global | us-west2-b | 2                  | Off                | dataproc-staging-us-west2-582004874449-adzhv9lz | Dec 11, 2020, 1:37:23 AM | Provisioning |

- Cluster Nodes with 1 master and 6 worker nodes

Google Cloud Platformbig-data-final

Search products and resources

Dataproc

Clusters

Jobs

Workflows

Autoscaling policies

Component exchange

Notebooks

pyspark-nlp-jupyter

SUBMIT JOBREFRESHDELETEVIEW LOGS

For PD-Standard without local SSDs, we strongly recommend provisioning 1TB or larger to ensure consistently high I/O performance. See <https://cloud.google.com/compute/docs/disks/performance> for information on disk I/O performance.

Name

Cluster UUID

Type

Status

pyspark-nlp-jupyter

051e89b8-b0d4-4a77-b428-78bd569134f5

Dataproc Cluster

Running

MONITORING

JOBS

VM INSTANCES

CONFIGURATION

WEB INTERFACES

Filter instances

|  | Name ↑                  | Role   |
|--|-------------------------|--------|
|  | pyspark-nlp-jupyter-m   | Master |
|  | pyspark-nlp-jupyter-w-0 | Worker |
|  | pyspark-nlp-jupyter-w-1 | Worker |
|  | pyspark-nlp-jupyter-w-2 | Worker |
|  | pyspark-nlp-jupyter-w-3 | Worker |
|  | pyspark-nlp-jupyter-w-4 | Worker |
|  | pyspark-nlp-jupyter-w-5 | Worker |

Equivalent [REST](#)

For the last task i.e. automating task 1 with airflow, an airflow cluster was set up using Google Cloud Composer that created and destroyed the dataproc (pyspark) cluster using an airflow scheduler.

Google Cloud Platform

big-data-final

Composer

Create environment

Name \*

pyspark-pipeline

## Node configuration

The configuration information for the Google Kubernetes Engine nodes running the Airflow software.

Node count \*

3

The number of nodes in the Google Kubernetes Engine cluster that will be used to run this environment.

Location \*

us-west2

The Google Compute Engine region where the environment will be created.

Zone

us-west2-b

The Compute Engine zone in which to deploy the VMs used to run the Apache Airflow software. If unspecified, the service will pick a zone in the Compute Engine region corresponding to the selected location. Must specify location before selecting zone. [Learn more.](#)

Machine type

n1-standard-1

The Google Compute Engine machine type used for cluster instances. If unspecified, the machine type will default to 'n1-standard-1'. Must specify location before selecting machine type. [Learn more.](#)

Disk size (GB)

20

The disk size in GB used for node VMs. Minimum is 20 GB. If unspecified, defaults to 100 GB. Cannot be updated.

OAuth Scopes

The set of Google API scopes to be made available on all node VMs. If empty, defaults to <https://www.googleapis.com/auth/cloud-platform>. Cannot be updated. [Learn more.](#)

Service account

582004874449-compute@developer.gserviceaccount.com

The Google Cloud Platform Service Account to be used by the node VMs. If a service account is not specified, the "default" Compute Engine service account is used. Cannot be updated.

- Next, we set up a kubernetes cluster using n1-standard v2 machines for master and 3 n1-standard v2 machines with a web server using n1-standard v2 machine to host Airflow GUI web page.

Google Cloud Platform

big-data-final

Composer

Create environment

Configuration of the Cloud SQL instance storing Airflow metadata

Cloud SQL machine type

db-n1-standard-2 (2 vCPU, 7.5 GB memory)

## Network configuration

The network configuration for the Google Kubernetes Engine cluster running the Airflow software.

☐ Enable VPC-native (using alias IP) ?

Network

The Google Compute Engine network ID to be used for machine communications. If unspecified, the network will be 'default'. Must specify location before selecting network.

Subnetwork

The Google Compute Engine subnetwork ID to be used for machine communications. Must specify network ID before selecting subnetwork ID.

Private IP ?

☐ Enable private IP

## Web server configuration

Configuration of the App Engine Flexible instance running Airflow web server

Web server network access control ?

☒ Allow access from all IP addresses

☐ Allow access only from specific IP addresses 

BETA

Web server machine type

composer-n1-webserver-2 (2 vCPU, 1.6 GB memory)



# Data Processing

To create a dataframe from the xml files, we are using python's [xml.etree.ElementTree](#) module.

**1a) Extracting user id and username and storing to the bucket:** [jupyter notebook link](#)

Raw Data:

```
: text_file = sc.textFile("gs://stackoverflow-dataset-677/Users.xml")
: text_file.take(3)
: ['<?xml version="1.0" encoding="utf-8"?>',
  '<users>',
  '  <row Id="-1" Reputation="1" CreationDate="2010-07-28T16:38:27.683" DisplayName="Community" LastAccessDate="2010-07-28T16:38:27.683" WebsiteUrl="http://meta.stackexchange.com/" Location="on the server farm" AboutMe="&lt;p&gt;Hi, I\'m not really a person.&lt;/p&gt;&#xA;&#xA;&lt;p&gt;I\'m a background process that helps keep this site clean!&lt;/p&gt;&#xA;&#xA;&lt;p&gt;I do things like&lt;/p&gt;&#xA;&#xA;&lt;ul&gt;&#xA;&lt;li&gt;Randomly poke old unanswered questions every hour so they get some attention&lt;/li&gt;&#xA;&lt;li&gt;Own community questions and answers so nobody gets unnecessary reputation from them&lt;/li&gt;&#xA;&lt;li&gt;Own downvotes on spam/evil posts that get permanently deleted&lt;/li&gt;&#xA;&lt;li&gt;Own suggested edits from anonymous users&lt;/li&gt;&#xA;&lt;li&gt;&lt;a href=&quot;http://meta.stackexchange.com/a/92006&quot;&gt;Remove abandoned questions&lt;/a&gt;&lt;/li&gt;&#xA;&lt;/ul&gt;&#xA;" Views="0" UpVotes="19522" DownVotes="185479" AccountId="-1" />']
```

- After cleaning, filtering and converting the data to dataframe, we get:

```
user_df.show(truncate=False)
```

| id | username         |
|----|------------------|
| 1  | Community        |
| 2  | Geoff Dalgas     |
| 3  | Jarrold Dixon    |
| 4  | txwikinger       |
| 5  | Nathan Osman     |
| 6  | Emmett           |
| 7  | Helix            |
| 8  | mechanical_meat  |
| 9  | Andrew           |
| 10 | DLH              |
| 11 | hannes.koller    |
| 12 | Michael Terry    |
| 13 | Keith Maurino    |
| 14 | Jweede           |
| 16 | Jeremy L         |
| 17 | tutuca           |
| 18 | excid3           |
| 20 | ParanoiaPuppy    |
| 21 | GeoD             |
| 22 | Alan Featherston |

only showing top 20 rows

- On completion of creating the dataframe, we store the data to GCS

```
In [82]: user_df1.count()
```

```
Out[82]: 855054
```

```
In [84]: user_df.repartition(1).write.csv("gs://stackoverflow-dataset-677/users_out", sep=',')
```

## 1b) Extracting the comments and joining username with Spark SQL: [jupyter notebook link](#)

Raw Data:

```
text_file.take(3)
```

```
['<?xml version="1.0" encoding="utf-8"?>',
 '<comments>',
 ' <row Id="13" PostId="23" Score="0" Text="Using /opt helps me keep track of the applications I\'ve installed myself." CreationDate="2010-07-28T19:36:59.773" UserId="10" ContentLicense="CC BY-SA 2.5" />']
```

- After cleaning, filtering and converting the data to dataframe, we get:

```
comments_df.show()
```

| postId | score | text                 | creationDate         | userId |
|--------|-------|----------------------|----------------------|--------|
| 23     | 0     | Using /opt helps ... | 2010-07-28T19:36:... | 10     |
| 18     | 0     | but popping in a ... | 2010-07-28T19:38:... | 10     |
| 27     | 0     | That will revert ... | 2010-07-28T19:39:... | 50     |
| 31     | 0     | I think you meant... | 2010-07-28T19:41:... | 12     |
| 18     | 0     | @DLH apparently n... | 2010-07-28T19:41:... | 63     |
| 12     | 2     | "ssh -X <server> ... | 2010-07-28T19:46:... | 96     |
| 12     | 0     | @Suppressingfire:... | 2010-07-28T19:48:... | 10     |
| 50     | 0     | Can you please re... | 2010-07-28T19:48:... | 56     |
| 27     | 0     | It probably shoul... | 2010-07-28T19:49:... | 5      |
| 58     | 0     | Do you mean the c... | 2010-07-28T19:50:... | 5      |
| 47     | 0     | Have you checked ... | 2010-07-28T19:50:... | 4      |
| 47     | 1     | Might be related ... | 2010-07-28T19:51:... | 104    |
| 58     | 0     | Do you use Gnome ... | 2010-07-28T19:51:... | 4      |
| 60     | 0     | This causes data ... | 2010-07-28T19:52:... | 66     |
| 18     | 0     | no the live CD do... | 2010-07-28T19:53:... | 4      |
| 52     | 0     | Does this let the... | 2010-07-28T19:55:... | 35     |
| 56     | 2     | LDAP and nfs are ... | 2010-07-28T19:56:... | 4      |
| 10     | 0     | Can I use it on a... | 2010-07-28T19:56:... | 27     |
| 70     | 1     | That's a good tip... | 2010-07-28T19:56:... | 45     |
| 70     | 0     | That is probably ... | 2010-07-28T19:58:... | 86     |

only showing top 20 rows

- Next, we read the username and user id from the csv file created in the previous step.



```
user_df.show()
```

```
+---+-----+
| id|      username|
+---+-----+
|  1|    Community|
|  2|   Geoff Dalgas|
|  3|   Jarrod Dixon|
|  4|    txwikinger|
|  5|   Nathan Osman|
|  6|      Emmett|
|  7|      Helix|
|  8| mechanical_meat|
|  9|      Andrew|
| 10|      DLH|
| 11| hannes.koller|
| 12| Michael Terry|
| 13| Keith Maurino|
| 14|      Jweede|
| 16|    Jeremy L|
| 17|      tutuca|
| 18|      excid3|
| 20| ParanoiaPuppy|
| 21|      GeoD|
| 22| Alan Featherston|
+---+-----+
only showing top 20 rows
```

- To allow join queries, we made the columns to be integer/long type.

```
comments_df.printSchema()
```

```
root
|-- postId: long (nullable = true)
|-- score: long (nullable = true)
|-- text: string (nullable = true)
|-- creationDate: string (nullable = true)
|-- userId: long (nullable = true)
```

```
user_df.printSchema()
```

```
root
|-- id: long (nullable = true)
|-- username: string (nullable = true)
```

- Converting both datasets to TempView for allowing SQL queries.

```
comments_df.createOrReplaceTempView("comments")
user_df.createOrReplaceTempView("users")
```

- Inner Join Query, introducing username to comments dataframe.

```
comments_users_sql_df = spark.sql("SELECT * FROM users u JOIN comments c ON u.id = c.UserId")
comments_users_sql_df.show()
```

|      | id  | username            | postId | score | text                 | creationDate         | userId |
|------|-----|---------------------|--------|-------|----------------------|----------------------|--------|
|      | 964 | Hendrik Brummermann | 4602   | 0     | I can confirm thi... | 2010-10-13T21:37:... | 964    |
|      | 964 | Hendrik Brummermann | 118087 | 0     | They took it in d... | 2012-04-28T06:17:... | 964    |
|      | 964 | Hendrik Brummermann | 638027 | 0     | I have the same i... | 2015-08-03T13:26:... | 964    |
| 1677 |     | eslambasha          | 84949  | 0     | @fossfreedom i do... | 2011-12-03T21:56:... | 1677   |
| 1697 |     | Frkstrem            | 16683  | 0     | @Marco, I know, I... | 2010-12-08T22:36:... | 1697   |
| 1697 |     | Frkstrem            | 16784  | 0     | This seems to be ... | 2010-12-09T19:05:... | 1697   |
| 1697 |     | Frkstrem            | 16886  | 1     | I only want to di... | 2010-12-10T22:26:... | 1697   |
| 1697 |     | Frkstrem            | 16892  | 1     | This is not an ac... | 2010-12-10T22:28:... | 1697   |
| 1697 |     | Frkstrem            | 16988  | 0     | Have you tried bu... | 2010-12-11T19:22:... | 1697   |
| 1697 |     | Frkstrem            | 17471  | 0     | @Stefano fixed it    | 2010-12-14T23:14:... | 1697   |
| 1697 |     | Frkstrem            | 17892  | 0     | My guess is that ... | 2010-12-17T13:50:... | 1697   |
| 1697 |     | Frkstrem            | 18014  | 0     | -1 It's too uncle... | 2010-12-18T17:53:... | 1697   |
| 1697 |     | Frkstrem            | 18273  | 0     | You did replace `... | 2010-12-22T17:48:... | 1697   |
| 1697 |     | Frkstrem            | 67121  | 0     | Firstly, I have a... | 2011-10-15T22:18:... | 1697   |
| 1697 |     | Frkstrem            | 108944 | 0     | You should use `t... | 2012-03-01T00:30:... | 1697   |
| 1697 |     | Frkstrem            | 453415 | 2     | Daily builds can ... | 2014-04-23T07:29:... | 1697   |
| 1697 |     | Frkstrem            | 223442 | 0     | @user2662639 Simp... | 2015-08-26T16:36:... | 1697   |
| 1697 |     | Frkstrem            | 223442 | 0     | @user2662639 (I t... | 2015-08-26T16:37:... | 1697   |
| 1697 |     | Frkstrem            | 17650  | 2     | @Fiksdal I don't ... | 2016-03-25T12:21:... | 1697   |
| 1697 |     | Frkstrem            | 899129 | 0     | @DavidFoerster Th... | 2017-04-01T13:36:... | 1697   |

## 2. Trending users with max comments:

- After extracting the comments data and converting to dataframe, we got:

|  | postId | score | text                 | creationDate         | userId |
|--|--------|-------|----------------------|----------------------|--------|
|  | 23     | 0     | Using /opt helps ... | 2010-07-28T19:36:... | 10     |
|  | 18     | 0     | but popping in a ... | 2010-07-28T19:38:... | 10     |
|  | 27     | 0     | That will revert ... | 2010-07-28T19:39:... | 50     |
|  | 31     | 0     | I think you meant... | 2010-07-28T19:41:... | 12     |
|  | 18     | 0     | @DLH apparently n... | 2010-07-28T19:41:... | 63     |
|  | 12     | 2     | "ssh -X <server> ... | 2010-07-28T19:46:... | 96     |
|  | 12     | 0     | @Suppressingfire:... | 2010-07-28T19:48:... | 10     |
|  | 50     | 0     | Can you please re... | 2010-07-28T19:48:... | 56     |
|  | 27     | 0     | It probably shoul... | 2010-07-28T19:49:... | 5      |
|  | 58     | 0     | Do you mean the c... | 2010-07-28T19:50:... | 5      |
|  | 47     | 0     | Have you checked ... | 2010-07-28T19:50:... | 4      |
|  | 47     | 1     | Might be related ... | 2010-07-28T19:51:... | 104    |
|  | 58     | 0     | Do you use Gnome ... | 2010-07-28T19:51:... | 4      |
|  | 60     | 0     | This causes data ... | 2010-07-28T19:52:... | 66     |
|  | 18     | 0     | no the live CD do... | 2010-07-28T19:53:... | 4      |
|  | 52     | 0     | Does this let the... | 2010-07-28T19:55:... | 35     |
|  | 56     | 2     | LDAP and nfs are ... | 2010-07-28T19:56:... | 4      |
|  | 10     | 0     | Can I use it on a... | 2010-07-28T19:56:... | 27     |
|  | 70     | 1     | That's a good tip... | 2010-07-28T19:56:... | 45     |
|  | 70     | 0     | That is probably ... | 2010-07-28T19:58:... | 86     |

- Next, we find the trending users by using `groupBy` aggregate function and sorting it with respect to count in descending order.

```
trending_users = comments_df.groupBy("userId") \
    .agg(count("text").alias("count")) \
    .where(col("userId") > 0) \
    .orderBy(col('count').desc())
```

- Result

```
trending_users.show()
```

```
+-----+-----+
|userId|count|
+-----+-----+
|167850|14677|
| 4272|12192|
|158442|12091|
| 15811|10505|
|175814| 8835|
|307523| 7861|
|469152| 7567|
|178692| 7523|
| 35795| 7515|
|126395| 7398|
|295286| 6757|
| 19421| 6679|
|344926| 6268|
| 19626| 6164|
| 22949| 6136|
| 10616| 6124|
|225694| 5736|
| 94914| 5692|
|459561| 5346|
| 72216| 5314|
+-----+-----+
```

- Next, we create a temp view of the trending users and merge it with usernames obtained using 1a.

```
In [17]: users_data = sc.textFile("gs://stackoverflow-dataset-677/users_out/*.csv")
```

```
In [18]: users_data.take(3)
```

```
Out[18]: ['1,Community', '2,Geoff Dalgas', '3,Jarrold Dixon']
```

```
In [19]: def create_user(rdd):
          rdd_split = rdd.split(",")
          return [int(rdd_split[0]),rdd_split[1]]
```

```
In [20]: users_rdd = users_data.map(lambda x: create_user(x))
```

```
In [21]: users_rdd.take(3)
```

```
Out[21]: [[1, 'Community'], [2, 'Geoff Dalgas'], [3, 'Jarrod Dixon']]
```

```
In [22]: user_data = ["id","username"]
          user_df = users_rdd.toDF(user_data)
```

- Joining the usernames with the trending user.


```
trending_usernames = spark.sql("SELECT count,username FROM trending_users tu JOIN users u ON u.id = tu.userId order by count desc limit 10")
trending_usernames.show()
```

| count | username         |
|-------|------------------|
| 14677 | Pilot6           |
| 12192 | heynnema         |
| 12091 | mur              |
| 10505 | Rinzwind         |
| 8835  | David Foerster   |
| 7861  | WinEunuuchs2Unix |
| 7567  | guiverc          |
| 7523  | steeldriver      |
| 7515  | Panther          |
| 7398  | oldfred          |


- Storing it to GCS, will be imported to Google BigQuery for visualization with Google Data Studio.

```
trending_usernames.repartition(1).write.csv("gs://stackoverflow-dataset-677/trending_usernames", sep=',')
```

- After importing the dataset to bigquery, and exporting the data to Google Data Studio for visualization.



- Result





### 3: Trending Topics in Comments:

[Notebook Link:](#)

- Obtaining the data from GCS and cleaning it.

```
In [3]: text_file = sc.textFile("hdfs:///pp-data/Comments.xml")

In [4]: text_file.take(3)

Out[4]: ['<?xml version="1.0" encoding="utf-8"?>',
 '<comments>',
 '  <row Id="13" PostId="23" Score="0" Text="Using /opt helps me keep track of the applications I\'ve installed myself." CreationDate="2010-07-28T19:36:59.773" UserId="10" ContentLicense="CC BY-SA 2.5" />']

In [5]: filteredRDD = text_file.filter(lambda x: x.startswith("  <row "))

In [6]: filteredRDD.take(1)

Out[6]: ['  <row Id="13" PostId="23" Score="0" Text="Using /opt helps me keep track of the applications I\'ve installed myself." CreationDate="2010-07-28T19:36:59.773" UserId="10" ContentLicense="CC BY-SA 2.5" />']

In [7]: cleanedRDD = filteredRDD.map(lambda x: x.lstrip(" "))

In [8]: cleanedRDD.take(1)

Out[8]: ['<row Id="13" PostId="23" Score="0" Text="Using /opt helps me keep track of the applications I\'ve installed myself." CreationDate="2010-07-28T19:36:59.773" UserId="10" ContentLicense="CC BY-SA 2.5" />']
```

- After extracting the attributes (comment text) from xml, we get

```
import xml.etree.ElementTree as ET

def parse_xml(rdd):
    """
    Read the xml string from rdd, parse and extract the elements,
    then return a list of list.
    """
    root = ET.fromstring(rdd)
    rec = []

    if "Text" in root.attrib:
        rec.append(root.attrib['Text'])
    else:
        rec.append("N/A")

    return rec
```

```
records_rdd = cleanedRDD.map(lambda x : parse_xml(x))
```

```
records_rdd.take(3)
```

```
[["Using /opt helps me keep track of the applications I've installed myself."],
 ["but popping in a live cd I already have isn't going to work huh?"],
 ['That will revert the splash screen as well as the login? I almost did that, but grew hesitant.']]
```

- Next, we have to remove the stop words from this text. So, we downloaded the stop words list online. Importing the stop words from file

```
stop_words_text = sc.textFile("file:///home/aarora7/P4-ayush-adarsh/03 stopwords.txt")

stop_words_text.take(3)

['a', 'about', 'above']
```

- Next, we converted it into dataframe. Converting the stop words dataframe to list
- 

```
: stop_words_list = list(stop_words_df.select('_1').toPandas()[ '_1'])
stop_words_list

: ['a',
  'about',
  'above',
  'across',
  'after',
  'again',
  'against',
  'all',
  'almost',
  'alone',
  'along',
  'already',
  'also',
  'although',
```

- Next, removing stopwords from main dataframe and converting into (stop\_word,1) tuple

```
: def remove_stop_word(x):
    list = []
    s_split = x[0].split(" ")
    for i in s_split:
        i = i.lower()
        if i not in stop_words_list:
            if not i.startswith('.') and not i.endswith('.') and not i.startswith('?') and not i.endswith('?') and not i.start
            swith('&') and not i.endswith('&'):
                list.append((i,1))
    return list

: cleaned_text = records_rdd.flatMap(lambda x: remove_stop_word(x))

: cleaned_text.take(3)

: [('/opt', 1), ('helps', 1), ('track', 1)]
```

- Next, finding the total count of words by using reduceByKey and converting it to Dataframe for storing to GCS.

```
: topic_count = cleaned_text.reduceByKey(lambda y,x: x+y)

: topic_count.take(3)

: [('manpage', 44), ('python', 6727), ('doesn't', 45536)]

: topic_count_col = ["topic", "count"]
topic_count_data = topic_count.toDF(topic_count_col)
```

- Storing to GCS and sorting the DF in descending order w.r.t count

```
final = topic_count_data.orderBy(col('count').desc())

final.repartition(1).write.csv("hdfs:///topic_count", sep=',')
```



- Next, we imported the data to BigQuery from GSC to export to Google Data Studio.

Google Cloud Platform big-data-final

BigQuery

Query editor

Query history

Saved queries

Job history

Transfers

Scheduled queries

Reservations

BI Engine

Resources + ADD DATA

Search for your tables and datasets

big-data-final-296808

trending\_username

topics

trending\_users


Run Save query Save view Schedule query More

topics


QUERY TABLE SHARE TABLE COPY TABLE DELETE TABLE

| Row | topic     | count  |
|-----|-----------|--------|
| 1   | Ubuntu    | 104944 |
| 2   | install   | 94016  |
| 3   | boot      | 51571  |
| 4   | sudo      | 45230  |
| 5   | version   | 39663  |
| 6   | Windows   | 37795  |
| 7   | partition | 28356  |
| 8   | apt-get   | 26273  |
| 9   | package   | 25401  |
| 10  | terminal  | 23624  |
| 11  | USB       | 22313  |
| 12  | kernel    | 20386  |
| 13  | windows   | 19935  |

- Visualizing the data using Google Data Studio.



- Result: Trending Keywords: Ubuntu is the most trending keyword with install as the second. This means most users comment about installation.



#### 4. Automating Part 1 with Google Cloud Composer (Apache Airflow).

- After setting up the server (shown in the beginning), next we create python scripts for part 1 and 2 and upload them to GCS. The GCS path will be referenced during the final orchestration.

##### Part 1: users.py

```
from pyspark.sql.functions import *
import time
import pyspark.sql.functions as F
from pyspark.sql.types import *
import xml.etree.ElementTree as ET
from pyspark.sql import SparkSession

spark = SparkSession \
    .builder \
    .appName("My PySpark code") \
    .getOrCreate()

text_file = sc.textFile("gs://stackoverflow-dataset-677/Users.xml")
filteredRDD = text_file.filter(lambda x: x.startswith(" <row "))
cleanedRDD = filteredRDD.map(lambda x: x.lstrip(" "))

def parse_xml(rdd):
    root = ET.fromstring(rdd)
    rec = []
```

```

    id = root.attrib['Id']
    if id == "-1":
        id = "1"
    rec.append(id)
    rec.append(root.attrib['DisplayName'])
    return rec

records_rdd = cleanedRDD.map(lambda x : parse_xml(x))

user_data = ["id", "username"]
user_df = records_rdd.toDF(user_data)
user_df.repartition(1).write.csv("gs://stackoverflow-dataset-677/users_out1", sep=',')

```

## Part 2: comments.py

```

from pyspark.sql.functions import *
import time
import pyspark.sql.functions as F
from pyspark.sql.types import *
import xml.etree.ElementTree as ET
from pyspark.sql import SparkSession

spark = SparkSession \
    .builder \
    .appName("My PySpark code") \
    .getOrCreate()

text_file = sc.textFile("gs://stackoverflow-dataset-677/Comments.xml")
filteredRDD = text_file.filter(lambda x: x.startswith(" <row "))
cleanedRDD = filteredRDD.map(lambda x: x.lstrip(" "))

def parse_xml(rdd):
    """
    Read the xml string from rdd, parse and extract the elements,
    then return a list of list.
    """
    root = ET.fromstring(rdd)
    rec = []

    if "PostId" in root.attrib:
        rec.append(int(root.attrib['PostId']))
    else:
        rec.append(0)

    if "Score" in root.attrib:
        rec.append(int(root.attrib['Score']))
    else:
        rec.append(0)

```

```

    if "Text" in root.attrib:
        rec.append(root.attrib['Text'])
    else:
        rec.append("N/A")
    if "CreationDate" in root.attrib:
        rec.append(root.attrib['CreationDate'])
    else:
        rec.append("N/A")
    if "UserId" in root.attrib:
        rec.append(int(root.attrib['UserId']))
    else:
        rec.append(0)
    return rec

records_rdd = cleanedRDD.map(lambda x : parse_xml(x))
comments_data = ["postId", "score", "text", "creationDate", "userId"]
comments_df = records_rdd.toDF(comments_data)
comments_df.createOrReplaceTempView("comments")
comments_sql_df = spark.sql("SELECT * FROM comments")


users_data = sc.textFile("gs://stackoverflow-dataset-677/users_out1/*.csv")

def create_user(rdd):
    rdd_split = rdd.split(",")
    return [int(rdd_split[0]), rdd_split[1]]

users_rdd = users_data.map(lambda x: create_user(x))
user_data = ["id", "username"]
user_df = users_rdd.toDF(user_data)
user_df.createOrReplaceTempView("users")
comments_users_sql_df = spark.sql("SELECT * FROM users u JOIN comments c ON u.id = c.UserId")
comments_users_sql_df.repartition(1).write.csv("gs://stackoverflow-dataset-677/users_out1",
sep=',')

```







- Next, we created two variables for the final orchestration. The GCP project id and the zone where the clusters will be created.


Airflow
DAGs
Data Profiling ▾
Browse ▾
Admin ▾
Docs ▾
About ▾

## Variables

No file chosen

List (2)
Create
Add Filter ▾
With selected ▾

| <input type="checkbox"/> |   | Key           | Val                   | Is Encrypted  |
|--------------------------|---|---------------|-----------------------|---|
| <input type="checkbox"/> |   | dataproc_zone | us-west2-b            |  |
| <input type="checkbox"/> |   | project_id    | big-data-final-296808 |  |

- **Final Orchestration:** To connect both the scripts together and for automatic creation and destruction of Spark Clusters, we created a DAG file.
  - This contains the path of the users.py and comments.py file and the airflow variables as well.

#### main\_dag.py

```
# Ref:
https://medium.com/analytics-vidhya/a-gentle-introduction-to-data-workflows-with-apache-airflow-and-apache-spark-6c2cd9aee573
from datetime import timedelta, datetime
from airflow import models
from airflow.operators.bash_operator import BashOperator
from airflow.contrib.operators import dataproc_operator
from airflow.utils import trigger_rule

# STEP 2: Define a start date
# In this case yesterday
yesterday = datetime(2020, 12, 10)

SPARK_CODE = ('gs://stackoverflow-dataset-677/01_user.py')
SPARK_CODE2 = ('gs://stackoverflow-dataset-677/02_user_comments_join.py')
dataproc_job_name = 'extract_users_job_dataproc'
dataproc_job_name2 = 'extract_comments_join_users_dataproc'

# STEP 3: Set default arguments for the DAG
default_dag_args = {
    'start_date': yesterday,
    'depends_on_past': False,
    'email_on_failure': False,
    'email_on_retry': False,
    'retries': 1,
    'retry_delay': timedelta(minutes=5),
    'project_id': models.Variable.get('project_id')
}

# STEP 4: Define DAG
# set the DAG name, add a DAG description, define the schedule interval and pass the
# default arguments defined before
with models.DAG(
    'comments_extract_user_join_spark_workflow',
    description='DAG for extracting comments and merging with user name',
    schedule_interval=timedelta(days=1),
    default_args=default_dag_args) as dag:

# STEP 5: Set Operators
# BashOperator
```

```

# A simple print date
print_date = BashOperator(
    task_id='print_date',
    bash_command='date'
)

# dataproc_operator
# Create small dataproc cluster
create_dataproc = dataproc_operator.DataprocClusterCreateOperator(
    task_id='create_dataproc',
    cluster_name='dataproc-cluster-demo-{{ ds_nodash }}',
    num_workers=2,
    zone=models.Variable.get('dataproc_zone'),
    master_machine_type='n1-standard-1',
    worker_machine_type='n1-standard-1')

run_spark = dataproc_operator.DataProcPySparkOperator(
    task_id='run_spark',
    main=SPARK_CODE,
    cluster_name='dataproc-cluster-demo-{{ ds_nodash }}',
    job_name=dataproc_job_name
)

run_spark2 = dataproc_operator.DataProcPySparkOperator(
    task_id='run_spark2',
    main=SPARK_CODE2,
    cluster_name='dataproc-cluster-demo-{{ ds_nodash }}',
    job_name=dataproc_job_name2
)

# dataproc_operator
# Delete Cloud Dataproc cluster.
delete_dataproc = dataproc_operator.DataprocClusterDeleteOperator(
    task_id='delete_dataproc',
    cluster_name='dataproc-cluster-demo-{{ ds_nodash }}',
    trigger_rule=trigger_rule.TriggerRule.ALL_DONE)

# STEP 6: Set DAGs dependencies
# Each task should run after have finished the task before.
print_date >> create_dataproc >> run_spark >> run_spark2 >> delete_dataproc

```



- To run the main dag file, we upload it to DAGs folder and the script ran automatically in 2 minutes.

The screenshot shows the Google Cloud Platform Storage interface. The left sidebar contains navigation options: Storage, Browser, Monitoring, and Settings. The main content area displays the 'Bucket details' for 'us-west2-pyspark-pipeline-770e6cc7-bucket'. Below the bucket name are tabs for OBJECTS, CONFIGURATION, PERMISSIONS, RETENTION, and LIFECYCLE. The 'OBJECTS' tab is active, showing a list of objects in the 'dags' folder. The objects are 'airflow\_monitoring.py' (729 B, text/x-python) and 'main\_dag.py' (2.1 KB, text/x-python-script). Both were created on Dec 9, 2020, and Dec 11, 2020, respectively.

| Name                  | Size   | Type                 | Created time              |
|-----------------------|--------|----------------------|---------------------------|
| airflow_monitoring.py | 729 B  | text/x-python        | Dec 9, 2020, 1:39:47 PM   |
| main_dag.py           | 2.1 KB | text/x-python-script | Dec 11, 2020, 12:09:18 AM |

- Main dag file with with workflow name “Comments\_extract\_user\_join\_spark\_workflow”.

The screenshot shows the Airflow DAGs list interface. The top navigation bar includes links for Airflow, DAGs, Data Profiling, Browse, Admin, Docs, and About. The main content area displays a table of DAGs. The table has columns for DAG, Schedule, Owner, Recent Tasks, Last Run, DAG Runs, and Links. Two DAGs are listed: 'airflow\_monitoring' and 'comments\_extract\_user\_join\_spark\_workflow'. The 'comments\_extract\_user\_join\_spark\_workflow' DAG is currently running, as indicated by the 'On' status and the '1 day, 0:00:00' schedule.

| DAG                                       | Schedule       | Owner   | Recent Tasks | Last Run         | DAG Runs | Links                    |
|---|----------------|---------|--------------|------------------|----------|--------------------------|
| airflow_monitoring                        | None           | airflow | 1 (1)        | 2020-12-11 09:37 | 1 (1)    | View, Edit, Delete, etc. |
| comments_extract_user_join_spark_workflow | 1 day, 0:00:00 | airflow | 1 (1)        | 2020-12-10 00:00 | 1 (1)    | View, Edit, Delete, etc. |

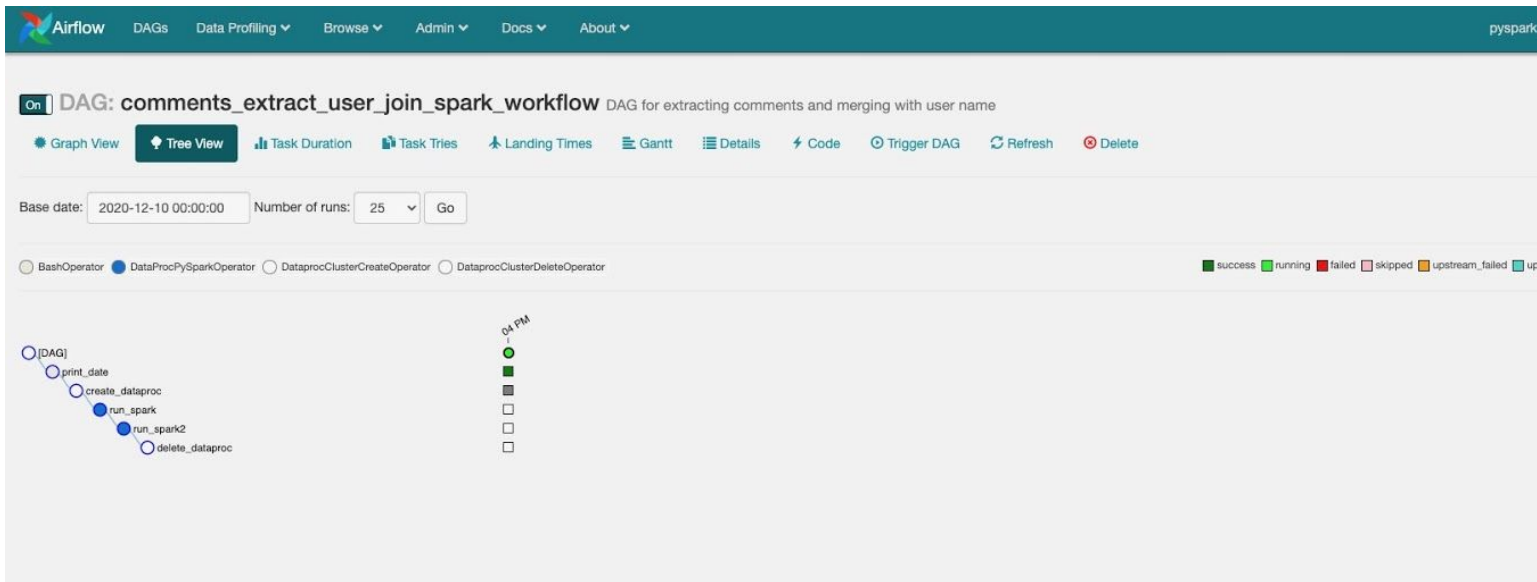
- Graph View of the Workflow

The screenshot shows the Airflow DAG Graph View for the 'comments\_extract\_user\_join\_spark\_workflow' DAG. The top navigation bar includes links for Airflow, DAGs, Data Profiling, Browse, Admin, Docs, and About. The main content area displays the DAG graph. The DAG is currently running, as indicated by the 'On' status and the 'running' label. The graph shows a sequence of tasks: 'print\_date', 'create\_dataproc', 'run\_spark', 'run\_spark2', and 'delete\_dataproc'. The 'run\_spark' task is highlighted in green, indicating it is the current task being executed.

```

graph LR
    print_date --> create_dataproc
    create_dataproc --> run_spark
    run_spark --> run_spark2
    run_spark2 --> delete_dataproc
  
```

- Tree View



- Finally, the output file was saved in GCS.